

Dimension Reduction Methods in Linear Regression

Yifei Sun

Contents

Principal components regression (PCR)	2
Partial least squares (PLS)	4
PCR and PLS using <code>caret</code>	5

```
library(ISLR)
library(pls)
library(caret)
```

Predict a baseball player's salary on the basis of various statistics associated with performance in the previous year. Use `?Hitters` for more details.

```
data(Hitters)
Hitters <- na.omit(Hitters)

set.seed(2021)
trRows <- createDataPartition(Hitters$Salary,
                               p = .75,
                               list = F)

# training data
x <- model.matrix(Salary ~ ., Hitters)[trRows, -1]
y <- Hitters$Salary[trRows]

# test data
x2 <- model.matrix(Salary ~ ., Hitters)[-trRows, -1]
y2 <- Hitters$Salary[-trRows]
```

Principal components regression (PCR)

We fit the PCR model using the function `pcr()`.


```
set.seed(2)
pcr.mod <- pcr(Salary ~ .,
               data = Hitters[trRows, ],
               scale = TRUE, # scale = FALSE by default
               validation = "CV")

summary(pcr.mod)
```

```
## Data:      X dimension: 200 19
## Y dimension: 200 1
## Fit method: svdpc
## Number of components considered: 19
##
## VALIDATION: RMSEP
## Cross-validated using 10 random segments.
##      (Intercept)  1 comps  2 comps  3 comps  4 comps  5 comps  6 comps
## CV              449    360.1   361.1   357.8   355.5   359.8   353.7
## adjCV           449    359.6   360.5   357.1   354.8   359.0   352.6
##      7 comps  8 comps  9 comps 10 comps 11 comps 12 comps 13 comps
## CV       354.8   357.2   358.5   364.8   369.3   369.6   373.2
## adjCV     353.6   355.9   357.1   363.0   367.2   367.7   370.9
##      14 comps 15 comps 16 comps 17 comps 18 comps 19 comps
## CV       369.6   370.9   361.9   369.1   364.5   367.0
## adjCV     366.9   368.1   359.1   365.9   361.2   363.6
```

```
##
## TRAINING: % variance explained
##      1 comps  2 comps  3 comps  4 comps  5 comps  6 comps  7 comps  8 comps
## X      40.20   60.14   71.10   78.94   84.21   88.82   92.35   95.07
## Salary  38.68   38.77   40.67   41.56   41.72   44.49   44.53   44.55
##      9 comps 10 comps 11 comps 12 comps 13 comps 14 comps 15 comps
## X      96.38   97.34   98.05   98.65   99.16   99.49   99.76
## Salary  44.86   45.21   45.34   45.37   46.04   48.71   49.01
##     16 comps 17 comps 18 comps 19 comps
## X      99.91   99.97   99.99   100.00
## Salary  51.87   51.87   53.35   53.46
```

```
validationplot(pcr.mod, val.type="MSEP", legendpos = "topright")
```



```
cv.mse <- RMSEP(pcr.mod)
ncomp.cv <- which.min(cv.mse$val[1,,]) - 1
ncomp.cv
```

```
## 6 comps
##      6
```

```
predy2.pcr <- predict(pcr.mod, newdata = Hitters[-trRows,],
                      ncomp = ncomp.cv)
# test MSE
mean((y2 - predy2.pcr)^2)
```

```
## [1] 103756.1
```

Partial least squares (PLS)


We fit the PLS model using the function `plsr()`.

```
set.seed(2)
pls.mod <- plsr(Salary~.,
  data = Hitters[trRows,],
  scale = TRUE,
  validation = "CV")

summary(pls.mod)
```

```
## Data:      X dimension: 200 19
## Y dimension: 200 1
## Fit method: kernelpls
## Number of components considered: 19
##
## VALIDATION: RMSEP
## Cross-validated using 10 random segments.
##      (Intercept)  1 comps  2 comps  3 comps  4 comps  5 comps  6 comps
## CV              449    356.9   356.2   357.7   367.4   372.1   373.9
## adjCV           449    356.3   355.0   356.4   365.6   369.2   370.1
##      7 comps  8 comps  9 comps 10 comps 11 comps 12 comps 13 comps
## CV          371.2   366.4   367.2   361.3   361.0   364.4   365.5
## adjCV       367.4   363.2   364.0   358.5   358.4   361.4   362.4
##      14 comps 15 comps 16 comps 17 comps 18 comps 19 comps
## CV          366.8   368.2   367.4   366.8   364.9   367.0
## adjCV       363.4   364.5   363.9   363.4   361.6   363.6
##
## TRAINING: % variance explained
##      1 comps  2 comps  3 comps  4 comps  5 comps  6 comps  7 comps  8 comps
## X          40.01   49.51   60.95   74.65   79.13   84.14   86.16   89.76
## Salary     40.75   44.71   45.91   46.68   49.00   50.39   51.39   51.76
##      9 comps 10 comps 11 comps 12 comps 13 comps 14 comps 15 comps
## X          92.77   94.43   96.71   97.78   98.31   98.64   99.06
## Salary     51.99   52.31   52.46   52.65   52.86   53.24   53.41
##      16 comps 17 comps 18 comps 19 comps
## X          99.43   99.93   99.95   100.00
## Salary     53.42   53.42   53.45   53.46
```

```
validationplot(pls.mod, val.type="MSEP", legendpos = "topright")
```



```
cv.mse <- RMSEP(pls.mod)
ncomp.cv <- which.min(cv.mse$val[1,,]) - 1
ncomp.cv
```

```
## 2 comps
##      2
```

```
predy2.pls <- predict(pls.mod, newdata = Hitters[-trRows,],
                      ncomp = ncomp.cv)
# test MSE
mean((y2 - predy2.pls)^2)
```

```
## [1] 104418.8
```

PCR and PLS using caret

PCR

```
ctrl1 <- trainControl(method = "cv",
                      selectionFunction = "best") # "oneSE" for the 1SE rule

# show information about the model
modelLookup("pcr")
```

```
##  model parameter      label forReg forClass probModel
## 1   pcr      ncomp #Components  TRUE   FALSE   FALSE
```

```
modelLookup("pls")
```

```
##  model parameter      label forReg forClass probModel
## 1   pls      ncomp #Components  TRUE   TRUE   TRUE
```

```
# Two ways for standardizing predictors
```

```
# train(..., preProc = c("center", "scale"))
set.seed(2)
pcr.fit <- train(x, y,
  method = "pcr",
  tuneGrid = data.frame(ncomp = 1:19),
  trControl = ctrl1,
  preProcess = c("center", "scale"))
```

```
predy2.pcr2 <- predict(pcr.fit, newdata = x2)
mean((y2 - predy2.pcr2)^2)
```


```
## [1] 103756.1
```

```
# pcr(..., scale = TRUE)
set.seed(2)
pcr.fit2 <- train(x, y,
  method = "pcr",
  tuneGrid = data.frame(ncomp = 1:19),
  trControl = ctrl1,
  scale = TRUE)
```

```
predy2.pcr3 <- predict(pcr.fit, newdata = x2)
mean((y2 - predy2.pcr3)^2)
```

```
## [1] 103756.1
```

```
ggplot(pcr.fit, highlight = TRUE) + theme_bw()
```




PLS

```
set.seed(2)
pls.fit <- train(x, y,
  method = "pls",
  tuneGrid = data.frame(ncomp = 1:19),
  trControl = ctrl1,
  preProcess = c("center", "scale"))
predy2.pls2 <- predict(pls.fit, newdata = x2)
mean((y2 - predy2.pls2)^2)
```

```
## [1] 107400.5
```

```
ggplot(pls.fit, highlight = TRUE)
```



Here are some old codes on ridge, lasso and ordinary least squares.

```
set.seed(2)
ridge.fit <- train(x, y,
  method = "glmnet",
  tuneGrid = expand.grid(alpha = 0,
    lambda = exp(seq(-1, 10, length=100))),
  trControl = ctrl1)
predy2.ridge <- predict(ridge.fit, newdata = x2)

set.seed(2)
lasso.fit <- train(x, y,
  method = "glmnet",
  tuneGrid = expand.grid(alpha = 1,
    lambda = exp(seq(-1, 5, length=100))),
  # preProc = c("center", "scale"),
  trControl = ctrl1)
predy2.lasso <- predict(lasso.fit, newdata = x2)
```

Comparing the models based on resampling results.

```
resamp <- resamples(list(lasso = lasso.fit,
  ridge = ridge.fit,
  pcr = pcr.fit,
```



```

                                pls = pls.fit))
summary(resamp)

##
## Call:
## summary.resamples(object = resamp)
##
## Models: lasso, ridge, pcr, pls
## Number of resamples: 10
##
## MAE
##           Min.   1st Qu.   Median     Mean   3rd Qu.     Max. NA's
## lasso 167.5977 208.5325 254.8819 241.2160 274.4900 299.1552    0
## ridge 175.1284 207.5289 230.4954 235.8461 255.7910 300.1394    0
## pcr   172.5466 200.4329 220.4705 230.5195 262.9214 297.0131    0
## pls   177.6785 207.5252 219.5236 232.5475 261.8803 301.0774    0
##
## RMSE
##           Min.   1st Qu.   Median     Mean   3rd Qu.     Max. NA's
## lasso 241.9932 267.6539 362.7664 346.0944 389.3080 540.4914    0
## ridge 258.8813 278.3543 314.6673 344.7360 387.0793 522.3878    0
## pcr   256.4354 280.3404 306.7773 341.6713 388.7912 537.0166    0
## pls   235.8924 281.6246 315.0572 344.2670 400.4012 535.7458    0
##
## Rsquared
##           Min.   1st Qu.   Median     Mean   3rd Qu.     Max. NA's
## lasso 0.04942920 0.2492379 0.4905530 0.4514729 0.6339149 0.7392623    0
## ridge 0.03183006 0.3390105 0.4625176 0.4518270 0.6157224 0.7482311    0
## pcr   0.03822160 0.3153999 0.4606597 0.4538285 0.6161780 0.7650883    0
## pls   0.02315483 0.4227511 0.5119100 0.4562174 0.5966268 0.6990407    0

bwplot(resamp, metric = "RMSE")

```

